NCSA's World Wide Web Server: Design and Performance*

Thomas T. Kwan[†] Robert E. McGrath[‡] Daniel A. Reed[§]

National Center for Supercomputing Applications Department of Computer Science

University of Illinois
Urbana, Illinois 61801

Abstract

The World Wide Web (WWW) server at the National Center for Supercomputing Applications (NCSA) is one of the most heavily accessed WWW servers in the world. This server is based on a collection of cooperating hosts that share a common file system. To increase our understanding of how users access this server and to provide a basis for assessing server and system software optimizations, we analyzed NCSA's server logs for multiple weeks during a five month period. This paper describes the server's architecture, presents the results of our access pattern analysis, and discusses the implications for server design and extensibility.

*An extended version of this paper has been submitted for possible publication in the proceedings of *Supercomputing '95*.

[†]Supported in part by the National Science Foundation and the Advanced Research Projects Agency under Cooperative Agreement NCR-8919038 with the Corporation for National Research Initiatives.

[‡]The National Center for Supercomputing Applications is funded by the National Science Foundation, the Advanced Research Projects Agency, corporate partners, and the state and University of Illinois.

[§]Supported in part by the National Science Foundation under grants NSF IRI 92-12976 and NSF CDA94-01124, by the National Aeronautics and Space Administration under Contract NAG-1-613, and by the Advanced Research Projects Agency under Contracts DABT63-91-C-0004 and DABT63-93-C-0040.

1 Introduction

The recent explosion of interest in the World Wide Web (WWW) [2] can be traced to the distribution of the CERN (European Laboratory for Particle Physics in Geneva, Switzerland) and NCSA (National Center for Supercomputing Applications) servers and WWW client browsers. In particular, NCSA Mosaic, the graphical user interface for WWW browsing, based on distributed, multimedia hypertext, has spawned several commercial variants and has made the Internet readily accessible to a much larger population than in the past (e.g., in late October 1994, the White House announced a WWW page).

Network statistics from Merit, the NSFNet backbone management group, show that WWW traffic is the largest and by far the fastest growing segment of the Internet, and growing numbers of government and commercial groups are making hundreds of gigabytes of data available via WWW servers. At the same time, the WWW servers at NCSA have experienced explosive growth in traffic, from one million requests per week in February 1994, to two million per week in June 1994, three million per week in September 1994, nearly four million per week in December 1994, and even larger numbers in 1995 [3].

To support continued growth, WWW servers must manage a multi-gigabyte (in some instances a multi-terabyte) data base of multimedia information, while concurrently serving multiple request streams. This places demands on the servers' underlying operating systems and file systems that lie far outside today's normal operating regime. Simply put, WWW servers must become more adaptive and intelligent. The first step on this path is understanding extant access patterns and responses. Based on this understanding, one can then develop more efficient and intelligent server and system file caching and prefetching strategies.

The remainder of this paper is organized as follows. In §2, we describe the architecture of NCSA's WWW server. This is followed in §3 by a description of the WWW data collected at NCSA. In §4–§5, we analyze this data to identify user request patterns and server responses. In §6, we discuss the implications of our findings. Finally, in §7, we conclude with a brief summary of our observations.

Begin Sidebar: The World Wide Web

The World Wide Web (WWW) is a global information system, providing hypertext-linked access to resources on the Internet. The WWW was developed at CERN, the European Particle Physics Laboratory in Geneva, Switzerland, and has spread around the world in a few years. Today, there are many thousands of information providers and millions of users on the WWW.

The WWW is based on a client-server architecture using its own protocols. The service is accessed through the Hypertext Transfer Protocol (HTTP), which allows any client to access any server. The Hypertext Markup Language (HTML), a subset of the Standard Generalized Markup Language (SGML), allows the specification of links from one document to another, and supports the construction of multimedia documents. The WWW also incorporates existing network services, such as FTP and Gopher.

Users access the WWW through one of many "browsers." The freely available NCSA Mosaic[1] browser, shown in Figure 1, has brought millions of users to the WWW, and spurred the development of new information sources and commercial browsers.

End Sidebar

2 NCSA WWW Server Architecture

The single most notable aspect of WWW traffic is its phenomenal growth rate. At NCSA, the traffic has grown from one million requests per week in February 1994 to nearly four million per week in December 1994 [3], with little prospect of a decline in the growth rate.

[1]NCSA Mosaic is a trademark of the Board of Trustees of the University of Illinois.

Shortly after NCSA's WWW server was established, it became clear that the volume of WWW traffic would stress operating systems and network implementations in ways not originally envisioned by their designers. At peak times, the NCSA server receives 30–40 new WWW requests per second, and because the Hypertext Transfer Protocol (HTTP) is connectionless, each such request appears to the server as a separate network connection.

Not only were most implementations of the TCP/IP network protocol not designed to accept connections at this sustained rate, even conservative projections of request rate growth showed that no single processor system could serve all requests. To support the growing request rate, NCSA has developed a scalable, WWW architecture that consists of a group of loosely coupled WWW servers. Though the servers operate independently, they collectively provide the illusion of a single server.

The server architecture is based on three components: (1) a collection of independent servers, (2) a WWW document tree shared among the servers and stored by the Andrew (AFS®)[2] [11] distributed file system, and (3) a round-robin domain name system (DNS) that multiplexes the domain name `www.ncsa.uiuc.edu` among the constituent servers.

With this architecture, the NCSA WWW *service* is always the same, although the number and identity of the particular *servers* may change from day to day. Beginning with one server in February 1994, the architecture grew to four servers in May, eight servers in November, and nine in early 1995. To meet increasing demand, NCSA will continue to add servers as needed.

Below, we briefly describe each of the three server components. For additional details on the server architecture, see Katz *et al* [4] and Kwan *et al* [5].

[2]The Andrew File System is a product of Transarc Corporation.

Begin Sidebar: Scalable Server Architecture

Development of the NCSA architecture required resolution of three key problems.

1. *Information addressing:* Externally, the NCSA server has a single domain name (`www.ncsa.uiuc.edu`). Incoming requests addressed to this domain name must be mapped to multiple servers, each with a separate, user-invisible domain name. This mapping allows NCSA to invisibly add servers to accommodate the growing number of incoming requests.

2. *Information distribution:* Each server must be capable of responding to requests for any portion of the NCSA WWW server data base. Otherwise, the servers must be more tightly coupled, an arbiter must distribute requests to servers based on request type, and it is likely that the arbiter will become the bottleneck.

3. *Load balancing:* The requests must be equally apportioned among the servers. Thus, newly added servers will always share the load and contribute to the scalability of the implementation.

End Sidebar

2.1 The Servers and the Network

The NCSA WWW server architecture is flexible enough to accommodate most Unix systems as component servers. The only requirements are that the systems function as Andrew File System (AFS) clients and support TCP/IP. The servers need not be homogeneous; the particular systems in use vary from time to time, and may be a heterogeneous collection of systems.

To date, the backbone of the NCSA WWW service has been a group of dedicated Hewlett-Packard HP 735 workstations. Though these systems are not generally considered "servers," their efficient TCP/IP implementation has made them an effective choice to process WWW requests. In the NCSA configuration, each HP 735 has 96 megabytes of memory and uses its local disk as a moderate size (130 megabytes) AFS cache. In addition, the local disk stores HTTP server log files and is the backing store for the virtual memory system.

The WWW servers are connected to the AFS file servers via a 100 megabit/second Fiber Distribution Data Interface (FDDI) ring; see Figure 2. The FDDI ring connects to the rest of NCSA and, via a T3 line, to the Internet.

2.2 NCSA AFS Configuration

All documents provided by the NCSA WWW service are served from the NCSA, center-wide Andrew File System (AFS) environment [7]. This distributed file system is shared by many hundreds of client workstations and supercomputers, as well as the WWW servers.

AFS provides a single, consistent view of the file system to each WWW server, allowing each server to access all of the WWW document tree. Because AFS clients (i.e., the WWW servers) cache recently used files on their local disks, the most frequently accessed documents are generally available locally, without remote disk access. In effect, AFS caching replicates the document tree on each WWW server.

Because AFS manages the shared document tree, the individual WWW servers need not and do not know either the number or identity of the other servers. It is difficult to over-emphasize the importance of this point. This allows rapid, "plug-and-play" addition (and removal) of component servers and the use of heterogeneous systems. In practice, we have found that servers can be added or removed from the ensemble in under an hour.

Our experience to date has been that AFS's local caching is critical to the success of the NCSA WWW server architecture. Stateless distributed file systems (e.g., NFS) cannot exploit the locality inherent in the HTTP request stream by locally caching frequently requested items. Instead, they must repeatedly retrieve those items from a shared file server. Not only does this increase the load on the file server, it is inherently unscalable.

Despite the advantages of local caching, much research remains to learn how well distributed file systems in general, and AFS in particular, support large, less frequently accessed files (e.g., 24-bit color images and digital video clips). With standard caching algorithms, access to these files will displace smaller, more frequently accessed files from the local cache. If these large, non-text files are not cached, access latencies for these items will be large. Data-type specific caching algorithms are one potential solution.

2.3 Round-Robin Domain Name System

The third and final component of the NCSA scalable WWW server is a modified network name resolver based on the Berkeley Internet Name Domain (BIND) code [1]. The existing BIND 4.9.2 code has a "round-robin" option that can associate a single domain name with several IP addresses. In response to requests, these addresses are distributed using a simple rotation algorithm. Because this rotation conflicted with extant software at NCSA, the BIND software was modified to only rotate specific addresses, namely those of the WWW servers; see [4] for details.

The modified domain name system (DNS) allows a domain name with more than one associated IP address to be specified as "round-robin." Each incoming request for the address of a round-robin domain name is satisfied with the next IP address on the list in a simple rotation. Thus, $\frac{1}{N}^{th}$ of the DNS requests get each of the N different IP addresses. This allows NCSA to maintain a group of WWW servers aliased by the single domain name `www.ncsa.uiuc.edu`. Adding a new server to the group is as simple as adding its IP address to the DNS entry for `www.ncsa.uiuc.edu`.

3 Data Collection

If the NCSA WWW service has accomplished nothing else, it has produced copious amounts of performance and access pattern data. These data are collected continuously on each server and are permanently archived each day to be available for researchers. Collectively, they constitute more than 150 megabytes of data each weekday [6]!

On each of the component WWW servers, the data collected includes

- the standard access logs from the NCSA HTTP daemons (httpd),
- the standard error logs from the httpd daemons,
- a custom log of the client browser type (the "user agent") that initiated each request,
- a trace of virtual memory statistics, obtained by recording UNIX *vmstat* data once each minute,
- a trace of packet counts, obtained by recording UNIX *netstat* data once each minute, and
- a count of active processes, sampled with *ps* once every five minutes.

Begin Sidebar: WWW Server Performance Visualization

To gain insights into the large volume of access and performance data in the WWW logs, we relied on a variety of standard statistical data analysis tools. However, to understand the dynamics of server behavior and the interactions of request patterns with round-robin DNS system, we exploited the local availability of the CAVE, an immersive, unencumbered virtual environment, and our Avatar visualization software [10], to create dynamic displays of server behavior. Figure 3 shows a snapshot of this visualization from a "day in the life" of the NCSA WWW server.

In the figure, the trajectories of four different servers in the performance metric space are denoted by the four colored ribbons. This snapshot, from near noon on September 7, 1994, shows that the round robin DNS system effectively balances the server load — the trajectories of all the servers cluster in the same region of the performance metric space; the small variations are due to differing request patterns.

End Sidebar

4 Request Pattern Analysis

To understand the access pattern and characteristics of NCSA's WWW service, we analyzed the data described in §3 for selected weeks during five different months of 1994. Below, we present the qualitative results with respect to the general access trends, the domain characteristics, and the file type distribution; see [5] for details.

4.1 General Trends

Qualitatively, the growth of WWW traffic on the Internet is well known. However, the specific characteristics of this growth and the sources of requests are much less well understood. Hence, the initial goal of our analysis was a simple characterization of traffic characteristics in terms of request count, request data volume, and request sources (by hardware platform type).

4.1.1 Traffic Growth

The number of requests received by the NCSA WWW servers during the period of our analysis has grown from about 300,000 per day in May to about 500,000 per day in September. Thus, the compounded growth rate over the five month period is roughly 14 percent per month. A scan of NCSA's January

1995 WWW server logs shows that the number of requests has increased to about 690,000 requests per day. As a result, the compounded growth rate is about 11 percent per month from May 1994 to January 1995. The slower growth rate may be due to the transient effect of the holiday season. A similar pattern was observed in January 1994.

4.1.2 Client Platforms

Knowing the platform from which a request originated has great potential value. Information providers can customize documents for different platforms, and servers can exploit this knowledge by tailoring their response to the hardware and software capabilities of the requesting platform.

The user agent logs from the first twenty days of December 1994 show that 31 percent of all connections were from X Windows clients, 38 percent from Microsoft Windows clients, 20 percent from Macintoshes, and 21 percent from all other types of clients. This data shows that at least 58 percent of the requests originate from personal computers. As vendors continue to ship new and improved versions of WWW browsers for personal computers, we expect requests from personal computers to grow at a very rapid rate. However, because of the relatively low bandwidth (modem) connections from most personal computers to the Internet, it is becoming increasingly important for WWW servers to adapt to client needs (e.g., by sending lower resolution images) and for clients to prefetch selected data to hide the long latency for data retrieval.

4.2 Domain Characteristics

Much discussion has centered on the commercial potential of the World Wide Web and the increasing accessibility of commercial information. To assess the number and distribution of commercial and other requesting sites, we aggregated domain names into a small number of broad categories: educational, commercial, government, and other. Table 1 summarizes the fraction of requests from the major Internet

Internet Domain	Percentage of Requests
Education (edu)	26
Commercial (com)	18
Government (gov)	5
Others	51

Table 1: Server Request Origins by Domain

domains.

Although Table 1 shows that the edu domain generates more requests than any other single domain, Figure 4 shows that the number of requests from commercial domains is growing rapidly. (For each month, the figure shows seven data points, corresponding to Sunday through Saturday of the week we analyzed during that month.) This reflects the increasing presence of commercial Internet service providers and the growing use of the Internet by the staff of commercial organizations.

Although the top ten educational and government domains (which generate the largest number of requests to NCSA's server) change almost daily, the top ten commercial domain names change little. Indeed, most of the top ten commercial domain names on any given day were also among the top ten domain names throughout the five months of data we analyzed.

The domain names in the com domain are mainly network firewalls for large organizations — they have long connection times and make unusually large number of requests. Because a firewall acts as a central location for accessing data outside a given organization, it is the ideal location for implementing network caching and proxy servers, a topic to which we return in §6.

4.3 Media Distributions

As we noted in §4.1.1, the request rate to the NCSA WWW server is growing at a compounded rate of between 11 and 14 percent per month. In addition to the rate, the characteristics of the growth have important implications for WWW server implementation. For example, satisfying large numbers of requests for small, text-based documents is much easier than responding to large numbers of requests for color images, video clips, or large data files.

Because the HTTPD server logs contain the name of the document being requested, and the file extension can be used to identify the document category, it is possible to determine the relative request frequency for text, images, audio, video, and data. The text category includes hypertext markup language (HTML) documents, plain text, and postscript files; the image category includes `GIF`, X bitmap (`xbm`), `JPEG`, and `RGB` files; the audio category includes `au`, `aiff`, and `aifc` files; and the video category includes `MPEG` and `QuickTime` files.

Figure 5 shows that text and images account for the majority of the requests. Although audio and video account for only one percent of the requests, they represent 28 percent of the bytes transferred. The requests for large audio and video files also lead to more bursty data transfer rates. Interestingly, the temporal distribution of the requests for audio and video is skewed toward later in the day than the distribution of text and images. We conjecture that users seek off-peak times to retrieve large items from the server.

One should be chary about projecting access characteristics from this data. The NCSA WWW document tree is dominated by a large number of small objects. As WWW document repositories mature, we expect them to contain a much larger number of large scientific and technical data sets, scientific visualizations and video clips, and audio segments. This shift will accentuate the behavior found in this study — many of the requests will be for small data items, but an increasing fraction of the data volume will be associated with requests for large, non-text items.

5 Server Caching

To this point, our focus has been on the characteristics of the request stream. We turn now to an examination of the servers' "response" to the incoming request stream.

Effective, distributed file caching was one of the key design principles in NCSA's WWW server architecture. Local caching at the WWW servers reduces the load on the shared AFS file servers, minimizes file traffic on the FDDI ring, and allows the WWW servers to respond quickly to requests for frequently accessed documents. To measure the effectiveness of the current AFS caching protocols, we analyzed the WWW server logs to identify the characteristics of the most frequently requested documents.

As mentioned in §2, NCSA serves documents from the AFS distributed file system, which automatically caches the most recently used files in local AFS client caches. Figure 6a shows the number of distinct files requested per day during the five months of our analysis; Figure 6b shows the total size of these same files.

Comparing Figures 6a and 6b shows that although the number of distinct files requested has increased, the total size of all the requested files has remained under 450 megabytes. Most of the newly added files have been small text and image files. To date, the AFS client cache hit ratios for the WWW servers have been near 90 percent, suggesting that AFS caching has worked quite well for the past access patterns.

Note that not only does the AFS file system cache frequently accessed files on the local disk of the WWW servers, the most frequently accessed of those files are cached in the primary memory of the WWW servers. With the observed access patterns to NCSA's WWW servers, less than 60 megabytes of primary memory cache space is needed to satisfy 95 percent of all incoming requests, which corresponds to roughly 800 distinct files. Though most requests are small, a small number of requests retrieve large items. For this reason, satisfying 95 percent of the requests represents only 80 percent of the total data volume.

6 Implications and Discussion

As the number of requests to NCSA's and other WWW servers continues to grow, the continued scalability of the server architecture, the efficiency of the HTTP protocol, and the effectiveness of caching strategies become increasingly critical research and implementation issues. Below, we discuss salient aspects of each issue.

6.1 Scalability and Persistent State

Although round robin DNS has allowed NCSA to add WWW servers without piercing the illusion that `www.ncsa.uiuc.edu` is a single host, the use of round robin DNS is not an ideal solution to either the decoupling of logical WWW server name from the physical server identity or to request load balancing. With this approach, the distribution of WWW server addresses is divorced from the characteristics and load of the constituent servers.

While the round robin mechanism equally distributes the IP addresses of the constituent servers, there is no mechanism to limit the number of times an address is used after it is distributed, nor to guarantee that the client system will honor the advertized time to live (TTL). For instance, a local DNS service might distribute a single IP address to any number of clients in its domain.

Moreover, envisioned extensions to HTTP include long lasting state (e.g., the results of previous data base searches) that must be retained by a WWW server. Supporting such extensions may be difficult for a multi-server architecture that relies on round robin DNS. A second request may be sent to a different server than the one holding the result of the previous request. Unless the data is shared (e.g., via AFS), obtaining the requisite information will require closer server cooperation, with associated overhead.

6.2 HTTP Protocol Extensions

The overriding trend from our data analysis is the continued growth in request rate. Currently, each request from the client uses a separate TCP connection, and the large number of short-lived TCP connections limits the performance of the server. This problem is exacerbated by the fact that a document may be composed of several pieces, each of which is fetched separately, with each fetch requiring a separate TCP/IP connection. Padmanabhan and Mogul [8] have proposed opening a single TCP connection per HTML document to avoid unnecessary TCP overhead; preliminary experiments show that this reduces document retrieval latency. Spero [12] has proposed a new protocol, HTTP-NG, which dramatically alters HTTP to reduce overhead, allow more parallelism, and efficiently support features such as authentication.

These and related protocol changes will reduce the latency to deliver data, and transmit more data over each TCP/IP connection. It will make HTTP servers much more like FTP and other session oriented services. This may well make much better use of the available network bandwidth and other server resources.

6.3 Distributed Caching and Prefetching

Beyond reducing the network protocol overhead, one can also aggressively cache and prefetch the data. At the moment, various browsers caches data on local client disks to improve performance. Pitkow and Recker [9] have shown that caching based on recent rates of past access is an effective technique. However, to design and implement effective prefetching, one must first study and understand the extant access patterns. Our data suggests that partitioned caches are a promising alternative. However, prototype implementations and trace-driven simulations are needed to measure the performance benefits that might accrue from this approach.

We noted in §4.2 that the most prolific sites are all commercial gateways. Moreover, about two percent of the requests to the NCSA WWW servers are from hosts that make only one request. The most popular of these requests are to the "directory" pages, namely the NCSA *Internet Starting Points*, the *Internet Resources Meta-Index*, and the *What's New* pages. These pages are excellent candidates for replication and caching throughout the Internet, particularly at commercial gateways.

In the future, as audio and video clips play a larger role in conveying multimedia information, audio and video requests will significantly affect network traffic and caching strategies. As we have seen, even a small increase in the use of these data types will dramatically increase the amount of data to be read and transmitted, with concomitant deleterious effect on the efficiency of server caching strategies.

7 Conclusions

We have described the design of NCSA's WWW server and analyzed the access patterns to the server in terms of the user request patterns and the responses of the server. The analysis shows that scalability, protocol efficiency, and effective caching strategies are the major issues for the next generation of WWW servers. In particular, we believe that both clients and servers must aggressively exploit caching and prefetching, based on knowledge of request patterns, data types, and hardware capabilities, to improve performance.

Acknowledgments

Our thanks to Eric Katz for providing us with the initial log analysis scripts, to Nancy Yeager, Michelle Butler, and Paul Zawada for providing crucial assistance in understanding the NCSA WWW server, and to Charlie Catlett, without whom this work would not have been possible. Finally, thanks to Will Scullin for developing the visualization software used to display dynamic server behavior.

References

[1] Albitz, P., and Liu, C. *DNS and BIND in a Nutshell.* O'Reilly and Associates, 1992.

[2] Berners-Lee, T., Cailliau, R., Luotonen, A., Nielsen, H., and Secret, A. The World-Wide Web. *Communications of the ACM 37*, 8 (Aug. 1994), 76–82.

[3] File, J. Web Server Activity: Total Connections, 1994. http://www.ncsa.uiuc.edu/SDG/ Presentations/Stats/WebServer.html.

[4] Katz, E. D., Butler, M., and McGrath, R. A Scalable HTTP Server: The NCSA Prototype. In *Proceedings of First International WWW Conference* (May 1994).

[5] Kwan, T. T., McGrath, R. E., and Reed, D. A. User Access Patterns to NCSA's World Wide Web Server. Tech. rep., University of Illinois at Urbana–Champaign, Department of Computer Science, February 1995.

[6] McGrath, R. E. What We Do and Don't Know About the Load on the NCSA WWW Server, Sept. 1994. http://www.ncsa.uiuc.edu/InformationServers/ Colloquia/28.Sep.94/Begin.html.

[7] NCSA. NCSA AFS Users Guide, 1994. http://www.ncsa.uiuc.edu/Pubs/UserGuides/ AFS-Guide/AFSv2.100.html.

[8] Padmanabhan, V. N., and Mogul, J. C. Improving HTTP Latency. In *Proceedings of Second International WWW Conference* (October 1994), pp. 995–1005.

[9] Pitkow, J. E., and Recker, M. M. A Simple Yet Robust Caching Algorithm Based on Dynamic Access Patterns. In *Proceedings of Second International WWW Conference* (October 1994), pp. 1039–1046.

[10] Reed, D. A., Shields, K. A., Tavera, L. F., Scullin, W. H., and Elford, C. L. Virtual Reality and Parallel Systems Performance Analysis. *IEEE Computer,* to appear (1995).

[11] Satyanarayanan, M. Scalable, Secure, and Highly Available Distributed File Access. *IEEE Computer 23*, 5 (May 1990), 9–21.

[12] Spero, S. Progress on HTTP-NG, 1994. http://www7.cern.ch/hypertext/WWW/Protocols/HTTP-NG/http-ng-status.html.

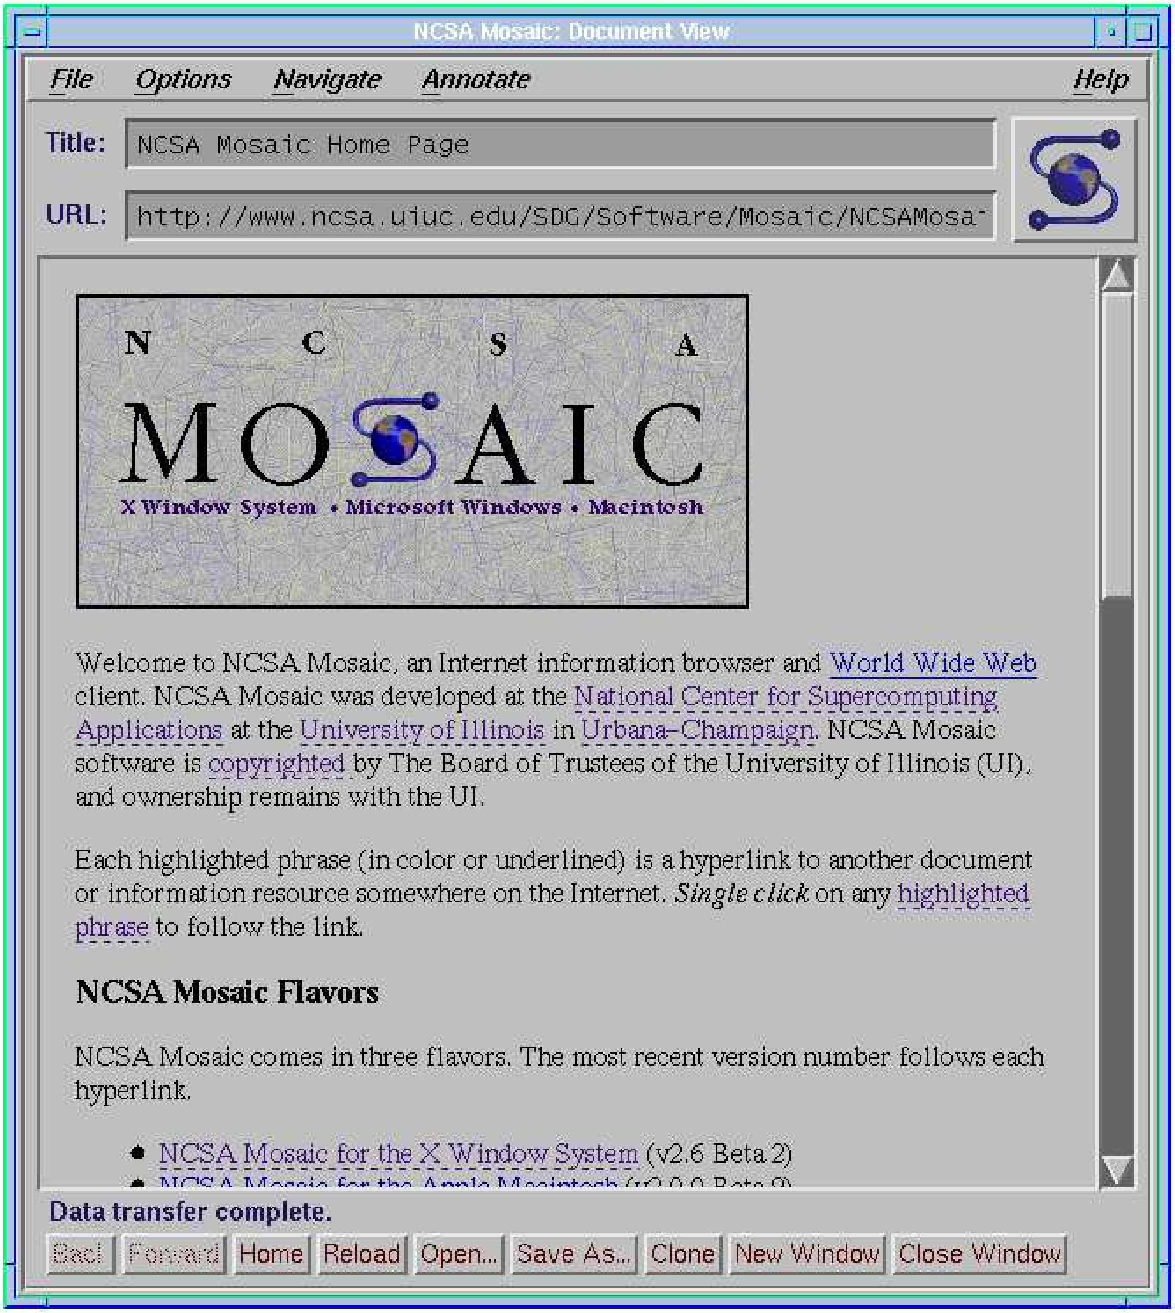

Figure 1: NCSA Mosaic WWW Browser

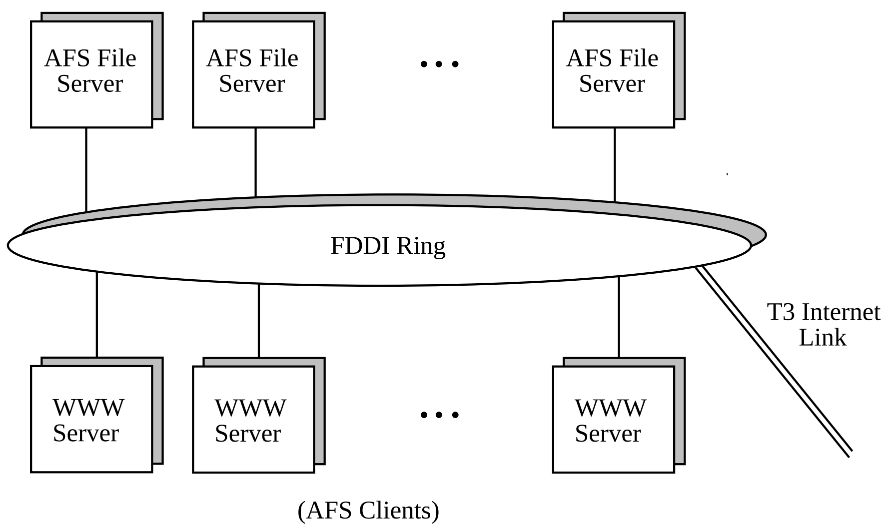

Figure 2: The NCSA Scalable WWW Server

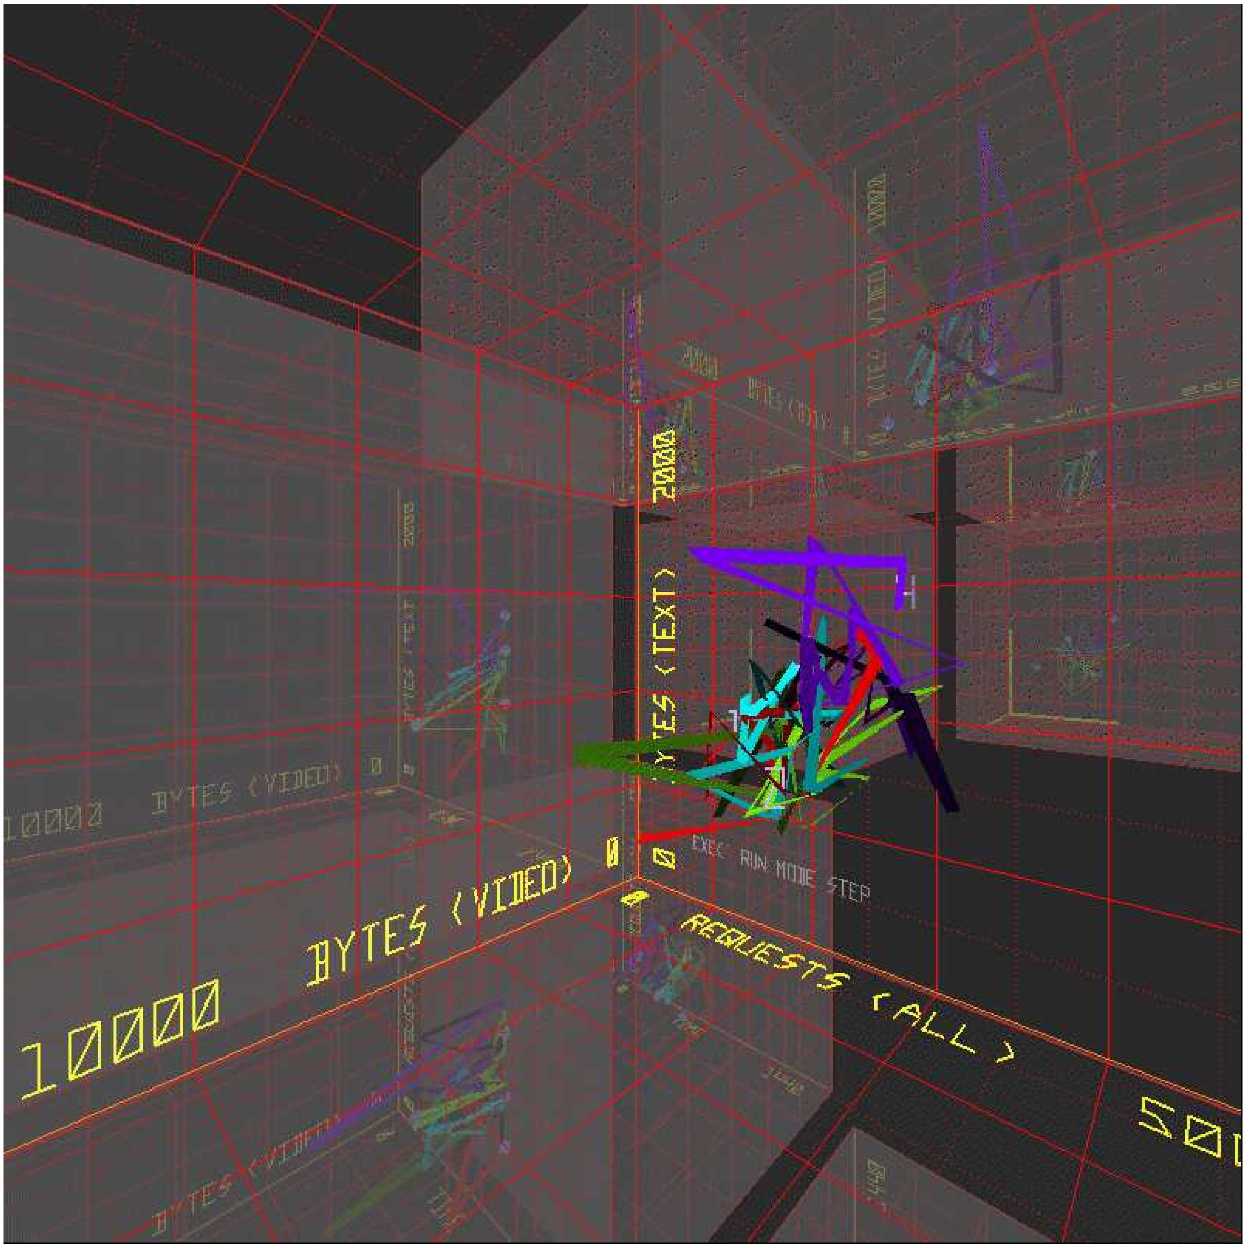

Figure 3: WWW Server Visualization

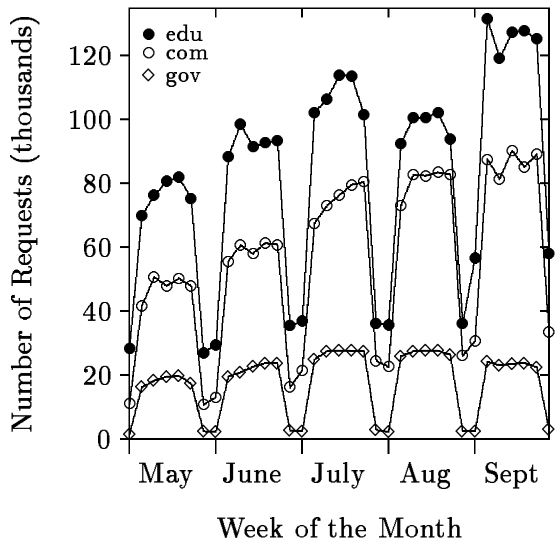

Figure 4: Weekly Domain Statistics

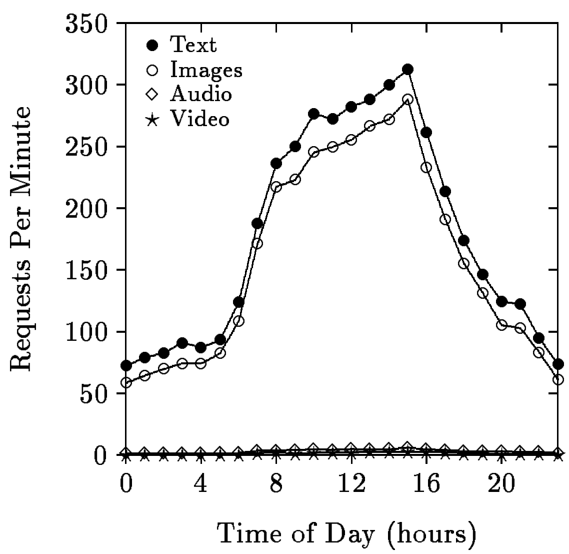

(a) Rate

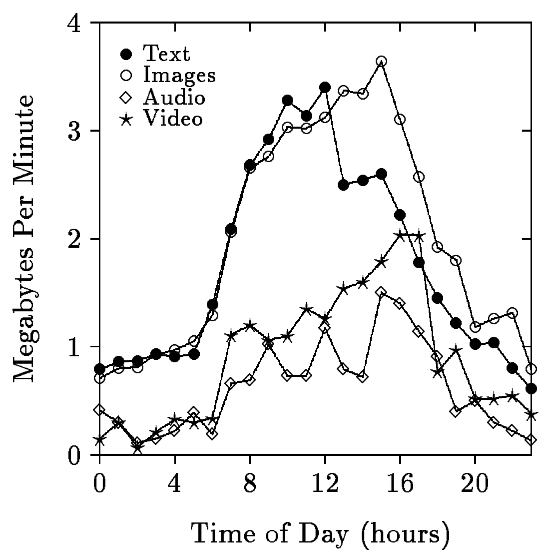

(b) Volume

Figure 5: File Type Statistics (September 7, 1994)

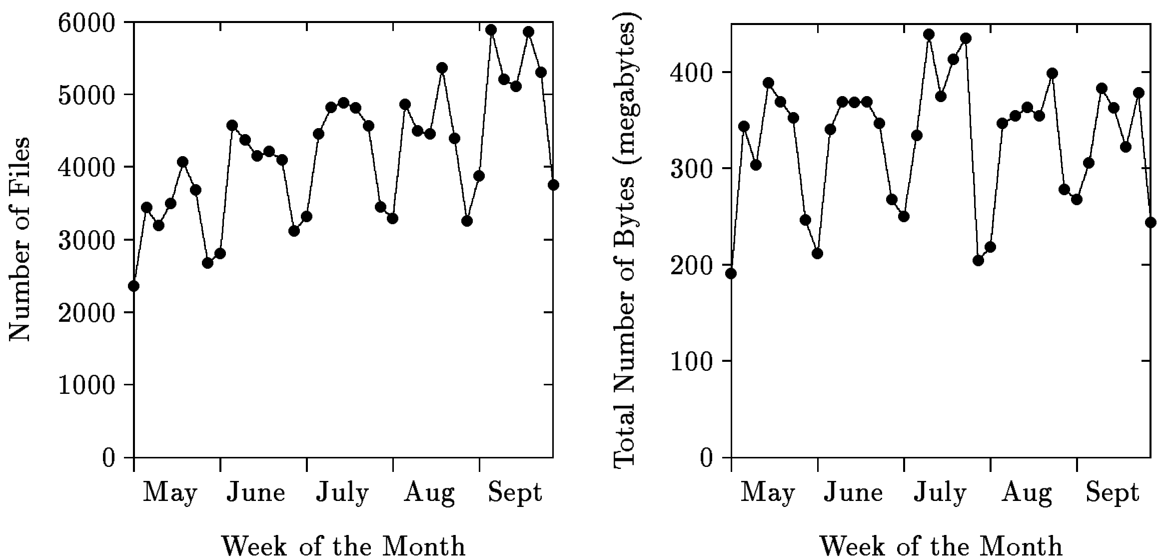

(a) Number of Distinct Files Requested

(b) Total Size of All Files Requested

Figure 6: Request Profile